
CLUSTERING

Sanjay Ranka
Distinguished Professor
Department of Computer and Information Science and Engineering
www.sanjayranka.com
sanjayranka@gmail.com
352 514 4213

1


DBSCAN

- DBSCAN is a density based clustering algorithm
 - Density = number of points within a specified radius (Eps)
 - A point is a *core point* if it has more than specified number of points ($MinPts$) within Eps
 - Core point is in the interior of a cluster
 - A *border point* has fewer than $MinPts$ within Eps but is in neighborhood of a core point
 - A *noise point* is any point that is neither a core point nor a border point
-

2

2


DBSCAN: Core, Border and Noise points



3

3

When DBSCAN works well




Original Dataset

Clusters found by DBSCAN


4

4

DBSCAN: Core, Border and Noise points



Original Points




Eps = 10, Minpts = 4
Point types:
Core
Border
Noise

5

5

DBSCAN: Determining Eps and MinPts


- Idea is that for points in a cluster, there k^{th} nearest neighbors are at roughly the same distance
- Noise points have the k^{th} nearest neighbor at a farther distance
- So, plot sorted distance of every point to its k^{th} nearest neighbor. ($k=4$ used for 2D points)




6

6


Where DBSCAN doesn't work well



Original Points



Minpts = 4, Eps = 9.75



MinPts = 4, Eps = 9.92

7

7


Subspace clustering

- Instead of using all the attributes (features) of a dataset, if we consider only subset of the features (subspace of the data), then the clusters that we find can be quite different from one subspace to another
- The clusters we find depend on the subset of the attributes that we consider

8

8

Subspace clustering




(a) Four clusters in three dimensions.

9

9

Subspace clustering




(b) View in the XY plane.

10

10

Subspace Clustering



(c) View in the XZ plane.

11

11


CLIQUE

- CLIQUE is a grid based clustering algorithm
- CLIQUE splits each dimension (attribute) in to a fixed number (ξ) of equal length intervals. This partitions the data space in to rectangular *units* of equal volume
- We can measure the density of each unit by the fraction of points it contains
- A unit is considered dense if its density $>$ user specified threshold T
- A cluster is a group of contiguous (touching) dense units

12

12

CLIQUE: Example



(a) Four clusters in three dimensions.




Figure 5.36. Histogram showing distribution of points for the X attribute.

13

13

CLIQUE

- CLIQUE starts by finding all the dense areas in the one dimensional spaces associated with each attribute
- Then it generates the set of two dimensional cells that might possibly be dense by looking at pairs of dense one dimensional cells
- In general, CLIQUE generates the possible set of k -dimensional cells that might possibly be dense by looking at dense $(k-1)$ -dimensional cells. This is similar to APRIORI algorithm for finding frequent item sets
- It then finds clusters by taking union of all adjacent high density cells

14

14

MAFIA

- Merging of Adaptive Finite Intervals (MAFIA) is a modification of CLIQUE that runs faster and finds better quality clusters.
- The main modification over CLIQUE is the use of an adaptive grid
- Initially each dimension is partitioned into a large number of intervals. A histogram is generated that shows the number of data points in each interval
- Groups of adjacent intervals are grouped in to windows, and the maximum number of points in the window's intervals becomes the value associated with the window

 15

15

MAFIA

- Adjacent windows are grouped together if the values of the two windows are close
- As a special case, if all windows are combined into one window, the dimensions is partitioned in to a fixed number of cells and the threshold for being considered a dense unit is increased for that dimension

 16

16

Limitations of CLIQUE and MAFIA

- Time complexity is exponential in the number of dimensions
- Will have difficulty if “too many” dense units are generated at lower stages
- May fail if clusters are of widely differing densities, since the threshold is fixed
- Determining the appropriate τ and ξ for a variety of data sets can be challenging
- It is not typically possible to find all clusters using the same threshold

17

17

Clustering Scalability for Large Datasets

- One very common solution is sampling, but the sampling could miss small clusters.
 - Data is sometimes not organized to make valid sampling easy or efficient.
- Another approach is to compress the data or portions of the data.
 - Any such approach must ensure that not too much information is lost.
(*Scaling Clustering Algorithms to Large Databases*, Bradley, Fayyad and Reina.)

18

18

Scalable Clustering: BIRCH

- BIRCH (Balanced and Iterative Reducing and Clustering using Hierarchies)
 - BIRCH can efficiently cluster data with a single pass and can improve that clustering in additional passes.
 - Can work with a number of different distance metrics.
 - BIRCH can also deal effectively with outliers.

19

19

Scaleable Clustering: BIRCH

- BIRCH is based on the notion of a clustering feature (CF) and a CF tree.
- A cluster of data points (vectors) can be represented by a triplet of numbers
 - (N, LS, SS)
 - N is the number of points in the cluster
 - LS is the linear sum of the points
 - SS is the sum of squares of the points.
- Points are processed incrementally.
 - Each point is placed in the leaf node corresponding to the “closest” cluster (CF).
 - Clusters (CFs) are updated.

20

20

Scalable Clustering: BIRCH

- Basic steps of BIRCH
 - Load the data into memory by creating a CF tree that “summarizes” the data.
 - Perform global clustering.
 - Produces a better clustering than the initial step.
 - An agglomerative, hierarchical technique was selected.
 - Redistribute the data points using the centroids of clusters discovered in the global clustering phase, and thus, discover a new (and hopefully better) set of clusters.

 21

21


Scalable Clustering: CURE

- Clustering Using Representatives
- Uses a number of points to represent a cluster
- Representative points are found by selecting a constant number of points from a cluster and then “shrinking” them toward the center of the cluster
- Cluster similarity is the similarity of the closest pair of representative points from different clusters
- Shrinking representative points toward the center helps avoid problems with noise and outliers
- CURE is better able to handle clusters of arbitrary shapes and sizes

 22

22


Experimental results: CURE



23

23


Experimental Results: CURE




24

24

CURE can not handle differing densities



Original Points



CURE

25

25

Graph Based Clustering

- Graph-Based clustering uses the proximity graph
 - Start with the proximity matrix
 - Consider each point as a node in a graph
 - Each edge between two nodes has a weight which is the proximity between the two points
 - Initially the proximity graph is fully connected
 - MIN (single-link) and MAX (complete-link) can be viewed as starting with this graph
- In the most simple case, clusters are connected components in the graph

26

26

Graph Based Clustering: Sparsification

- The amount of data that needs to be processed is drastically reduced
 - Sparsification can eliminate more than 99% of the entries in a similarity matrix
 - The amount of time required to cluster the data is drastically reduced
 - The size of the problems that can be handled is increased

27

27


Sparsification

- Clustering may work better
 - Sparsification techniques keep the connections to the most similar (nearest) neighbors of a point while breaking the connections to less similar points.
 - The nearest neighbors of a point tend to belong to the same class as the point itself.
 - This reduces the impact of noise and outliers and sharpens the distinction between clusters.
- Sparsification facilitates the use of graph partitioning algorithms (or algorithms based on graph partitioning algorithms).
 - Chameleon and Hypergraph-based Clustering

28

28

Sparsification




29

29

Limitations of Current Merging Schemes

- Existing merging schemes are static in nature



30

30

Chameleon: Clustering Using Dynamic Modeling

- Adapt to the characteristics of the data set to find the natural clusters.
- Use a dynamic model to measure the similarity between clusters.
 - Main property is the relative closeness and relative inter-connectivity of the cluster.
 - Two clusters are combined if the resulting cluster shares certain *properties* with the constituent clusters.
 - The merging scheme preserves *self-similarity*.

- One of the areas of research is the clustering of spatial data. a.





31

31

Characteristics of Spatial Datasets

- Clusters are defined as densely populated regions of the space
- Clusters have arbitrary shapes, orientation, and non-uniform sizes
- Difference in densities across clusters and variation in density within clusters
- Existence of special artifacts (*streaks*) and noise
- The clustering algorithm must address the above characteristics and also require minimal supervision



32

32

Chameleon

- Preprocessing Step: Represent the Data by a Graph
 - Given a set of points, we construct the *k-nearest-neighbor* (*k-NN*) graph to capture the relationship between a point and its *k* nearest neighbors.
- Phase 1: Use a multilevel graph partitioning algorithm on the graph to find a large number of clusters of well-connected vertices.
 - Each cluster should contain mostly points from one “true” cluster, i.e., is a sub-cluster of a “real” cluster.
 - Graph algorithms take into account global structure.

33

33


Chameleon

- Phase 2: Use Hierarchical Agglomerative Clustering to merge sub-clusters.
 - Two clusters are combined if the resulting cluster shares certain *properties* with the constituent clusters.
 - Two key properties are used to model cluster similarity:
 - Relative Interconnectivity: Absolute interconnectivity of two clusters normalized by the internal connectivity of the clusters.
 - Relative Closeness: Absolute closeness of two clusters normalized by the internal closeness of the clusters.

34

34


Experimental Results: Chameleon



35

35


Experimental Results: Chameleon



36

36


Experimental Results: CURE (10 clusters)



37

37


Experimental Results: CURE (15 clusters)



38

38


Experimental Results: Chameleon



39

39


Experimental Results: CURE (9 clusters)



40

40

Experimental Results: CURE (15 clusters)



41